

Phase 0: Foundations

Status: In Progress **Started:** 2024-11-11 **Target Duration:** 4 weeks (hardcore mode)

Philosophy: No checkbox learning. Every exercise must force you to think deeply about memory. If you can complete something without understanding what's happening at the memory level, it's not hard enough.

Learning Objectives

By the end of this phase, you must be able to:

- Read and write C code involving complex pointer manipulations
- Debug segfaults and memory issues using proper tools
- Explain what happens at the memory/syscall level for malloc/free
- Read basic x86-64 assembly generated from C code
- Implement a simple memory allocator from scratch
- Use gdb, valgrind, and AddressSanitizer fluently

Tools Setup (Week 1, Day 1 - NON-NEGOTIABLE)

Required Tools

```
# Compiler with debug support
gcc --version # or clang

# Debugger
gdb --version # or lldb on Mac

# Memory debugging
valgrind --version # Use Docker on Mac if needed

# Compilation flags you'll use CONSTANTLY
gcc -Wall -Wextra -g -O0 -fsanitize=address program.c
```

Setup Checklist

- Install gcc/clang with debug symbols
- Install gdb (Linux) or lldb (Mac)
- Install valgrind (or setup Docker container for Mac)
- Test AddressSanitizer: compile with -fsanitize=address
- Install strace (Linux) or dtraceunless (Mac)
- Setup godbolt.org account for assembly viewing

Tool Practice (Day 1)

- Create intentional segfault: `int *p = NULL; *p = 42;`
- Debug with gdb: set breakpoint, step, print variables
- Detect with valgrind: `valgrind ./program`
- Catch with ASan: compile with `-fsanitize=address`
- View syscalls: `strace ./program` (Linux) or `dtruss ./program` (Mac)

Resources

Primary (READ THESE)

- "The C Programming Language" (K&R) - Chapters 1-5, 6, 7
- "Computer Systems: A Programmer's Perspective" (CS:APP)
 - Chapter 3: Machine-Level Representation (x86-64 assembly)
 - Chapter 9: Virtual Memory (sections 9.1-9.9)
- "What Every Programmer Should Know About Memory" (Ulrich Drepper) - First 40 pages

Supplementary

- CS50 lectures on C and memory
- "Beej's Guide to C Programming"
- "Understanding glibc malloc" (online article)
- Source code: glibc malloc.c (simplified version)

Online Tools

- godbolt.org - Compiler Explorer (C to Assembly)
- cdecl.org - Decode complex C declarations
- onlinegdb.com - Quick debugging

Week-by-Week Breakdown

Week 1: Strings, Pointers, and Tools

Goal: Master pointer basics and tool usage. No more segfaults you can't debug.

Exercises

- **String functions from scratch** (NO <string.h>):
 - my_strlen(const char *s)
 - my_strdup(char *dst, const char *src)
 - my_strcat(char *dst, const char *src)
 - my_strcmp(const char *s1, const char *s2)
- **Understand memory segments:**

```
char str1[] = "hello";           // Where is this?  
char *str2 = "hello";           // Where is this?  
char *str3 = malloc(6);         // Where is this?  
  
// Write program that prints addresses and explains
```

- **Pointer arithmetic mastery:**
 - Array traversal using pointers (no [])
 - Reverse string in-place using two pointers
 - Implement void *my_memcpy(void *dst, const void *src, size_t n)
- **Tool practice:**
 - Deliberately create 5 different types of bugs
 - Debug each with gdb
 - Detect each with valgrind
 - Fix and verify

Reading

- K&R Chapters 1-5
- CS:APP 3.1-3.5 (Assembly basics)

Deliverable

- week1/ folder with all exercises
- Each .c file must compile with -Wall -Wextra -Werror
- README explaining what you learned about memory layout

Week 2: Assembly and Memory Layout

Goal: Understand what the CPU actually executes. No more abstractions.

Exercises

- **C to Assembly exploration:**
 - Write simple functions, compile with `gcc -S`
 - Understand basic instructions: `mov`, `add`, `call`, `ret`
 - See how variables map to registers/stack
 - Use godbolt.org to see optimizations (-O0 vs -O2 vs -O3)
- **Stack deep dive:**
 - Write recursive function, view stack frames in `gdb`
 - Implement `void inspect_stack(void)` that prints stack addresses
 - Understand function prologue/epilogue
 - Deliberately overflow stack, see what happens
- **Heap exploration:**
 - `malloc()` 10 times, print addresses - see the pattern?
 - Use `strace` to see `brk()`/`mmap()` syscalls
 - Free in different orders, observe behavior
 - Double-free detection with ASan
- **Reading real code:**
 - Clone GNU coreutils
 - Read and annotate `cat.c` source
 - Read and annotate `ls.c` source (focus on memory management)

Reading

- CS:APP Chapter 3.6-3.10 (Assembly continued)
- CS:APP Chapter 9.1-9.5 (Virtual Memory basics)
- "Understanding glibc malloc" article

Deliverable

- `week2/` folder with:
 - Annotated assembly examples
 - Stack/heap exploration program
 - Notes on `cat.c/ls.c` code reading
 - Answers to malloc investigation (see `questions.md`)

Week 3: Data Structures and Memory Management

Goal: Build complex structures, manage memory correctly, no leaks.

Exercises

- **Dynamic array (vector):**

```
typedef struct {  
    int *data;  
    size_t size;  
    size_t capacity;  
} Vector;  
  
// Implement:  
Vector* vec_create(void);  
void vec_push(Vector *v, int value);  
int vec_get(Vector *v, size_t index);  
void vec_destroy(Vector *v);
```

- Must resize when full (capacity doubling)
 - Zero leaks with valgrind
 - Write tests for edge cases
- **Linked list (from scratch):**

```
typedef struct Node {  
    int data;  
    struct Node *next;  
} Node;  
  
// Implement:  
Node* list_create(int data);  
void list_append(Node **head, int data);  
void list_delete(Node **head, int data);  
void list_destroy(Node **head);  
Node* list_reverse(Node *head);
```

- Handle empty list, single node
 - Implement reverse IN-PLACE (no extra memory)
 - Zero leaks
- **Pointer to pointer mastery:**

- Understand why `list_append(Node **head, ...)` needs `**`
 - Implement tree insertion using pointer-to-pointer
 - Draw memory diagrams
- **Memory leak debugging:**

- Take a provided buggy program (create one with 10 leaks)
- Find all leaks with valgrind
- Fix systematically
- Verify with valgrind --leak-check=full

Reading

- K&R Chapter 6 (Structures)
- CS:APP 9.9 (Dynamic Memory Allocation)

Deliverable

- week3/ folder with:
 - Vector implementation + tests
 - Linked list implementation + tests
 - Leak debugging report
 - Memory diagrams for pointer-to-pointer

Week 4: Mini-Project - Memory Allocator

Goal: Implement your own malloc/free. Understand heap management from the inside.

Project: my_malloc / my_free

Specifications:

```
void* my_malloc(size_t size);
void my_free(void *ptr);
```

Requirements:

- Use sbrk() (simple) or mmap() (better) to get memory from OS
- Maintain a free list of available blocks
- Implement first-fit or best-fit allocation strategy
- Store metadata (block size, free/used flag) BEFORE each block
- Implement coalescing of adjacent free blocks
- Handle alignment (8-byte or 16-byte boundaries)
- Handle edge cases: size=0, free(NULL), double-free

Implementation phases:

1. **Day 1-2:** Basic version - just sbrk() and linear search
2. **Day 3-4:** Add free list and coalescing
3. **Day 5:** Add alignment and metadata
4. **Day 6-7:** Testing, debugging, benchmarking

Testing:

- Allocate → Free → Allocate (verify memory reuse)
- Allocate many small blocks, free in different orders
- Stress test: 10000 random malloc/free operations
- No leaks with valgrind
- Benchmark vs system malloc (don't worry if slower)

Resources:

- Tutorial: "A malloc Tutorial" (online)
- Read simplified dlmalloc source code
- CS:APP 9.9.12 (Explicit Free Lists)

Deliverable

- week4/my_allocator/ folder with:
 - my_malloc.c / my_malloc.h
 - test_allocator.c with comprehensive tests
 - benchmark.c comparing to system malloc
 - DESIGN.md explaining your design choices
 - Valgrind report showing zero leaks
-

Key Concepts (Fill as you learn)

Pointers

What I've learned:

- [Date] ...
- [Date] ...

Stack vs Heap

What I've learned:

- [Date] ...

Memory Layout (Text, Data, BSS, Heap, Stack)

What I've learned:

- [Date] ...

Assembly Basics

What I've learned:

- [Date] ...

Memory Allocation Strategies

What I've learned:

- [Date] ...
-

Reality Checks

Week 1 Check (Date: _____)

- Can I debug a segfault in under 5 minutes?
- Do I understand why `char *s = "hello"; s[0] = 'H';` crashes?
- Can I explain pointer arithmetic to someone?

Week 2 Check (Date: _____)

- Can I read basic x86-64 assembly?
- Do I understand calling conventions?
- Can I explain what malloc() does at the syscall level?

Week 3 Check (Date: _____)

- Can I implement any data structure without memory leaks?
- Do I understand pointer-to-pointer usage?
- Can I draw memory layout diagrams?

Week 4 Check (Date: _____)

- Did I actually implement malloc, or just copy code?
 - Can I explain heap fragmentation?
 - Am I ready for Phase 1, or do I need more time?
-

Milestones

- Compiled first program with `-Wall -Wextra -Werror` with no warnings
- Debugged first segfault using gdb

- Found first memory leak with valgrind
 - Implemented linked list with zero leaks
 - Read first 100 lines of real C codebase
 - Read first assembly output from compiled C
 - Completed working memory allocator
 - **Phase 0 complete - READY FOR ASSEMBLY & MEMORY SYSTEMS**
-

Success Criteria (ALL must be true)

Before moving to Phase 1, you MUST:

1. Implement working my_malloc/my_free with zero leaks
2. Be able to debug segfaults without help
3. Read basic x86-64 assembly from compiled C
4. Understand stack, heap, data, text segments
5. Have zero fear of pointers, including pointer-to-pointer

If any is false, extend Phase 0. No shortcuts.

Started: 2024-11-11 | Target Completion: 2024-12-09 | Actual: TBD